TaD: A Plug-and-Play Task-Aware Decoding Method to Better Adapt LLMs on Downstream Tasks

Xinhao Xu, Hui Chen, Zijia Lin, Jungong Han, Lixing Gong,
Guoxin Wang, Yongjun Bao, Guiguang Ding

IJCAI JEJU 2024

Introduction

- **Fine-tuning as a common strategy to enhance the pre-trained LLMs in downstream tasks:**
  - Algorithmic side: better fine-tuning methods, e.g. PEFT.
  - Data side: more effective datasets.
  - Inherent knowledge acquisition of fine-tuned LLMs **rarely** investigated in the existing works.
- **Motivation:** The outputs of pre-trained LLMs do not always accurately reflect the knowledge they possess.

Research Problem

How can we leverage such inherent knowledge in the fine-tuned LLMs to enhance their performance in downstream tasks?

- **Intuitive Ideas:**
  - Token-predicting behavior alterations during the fine-tuning process reflect the the inherent knowledge.
  - Such alterations indicate an adaptive shift from common knowledge to specific knowledge for downstream tasks.
  - Manually mining and leveraging such inherent knowledge can improve the adaptation of LLMs on downstream tasks.

Method

- **A demonstration of the proposed knowledge vector and TaD:**

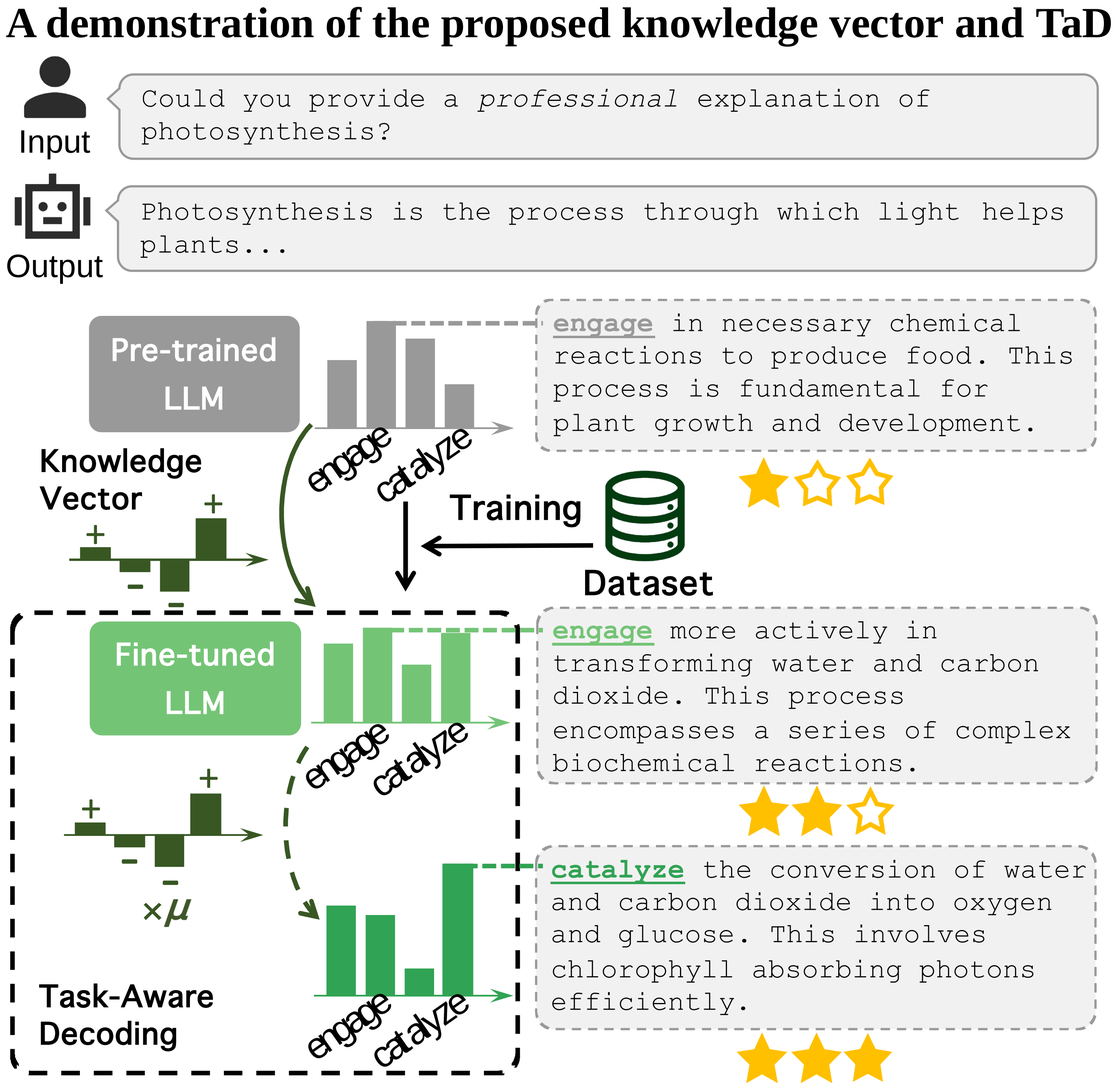

- **Knowledge Vector:**
  - Formulating the knowledge difference.
  - Explicitly denoting the direction of knowledge adaptation learned by a pre-trained LLM during fine-tuning.
  - Naturally possessing semantic information.
- **Task-Aware Decoding:**
  - Enhancing the fine-tuned LLM's output probability distribution with the knowledge vector.
  - Reinforcing the model's knowledge adaptation to downstream tasks for better performance.

- **A simplified illustration of our work:**
  - p_θ and p_ϕ: predicted token distribution of pre-trained and fine-tuned LLMs

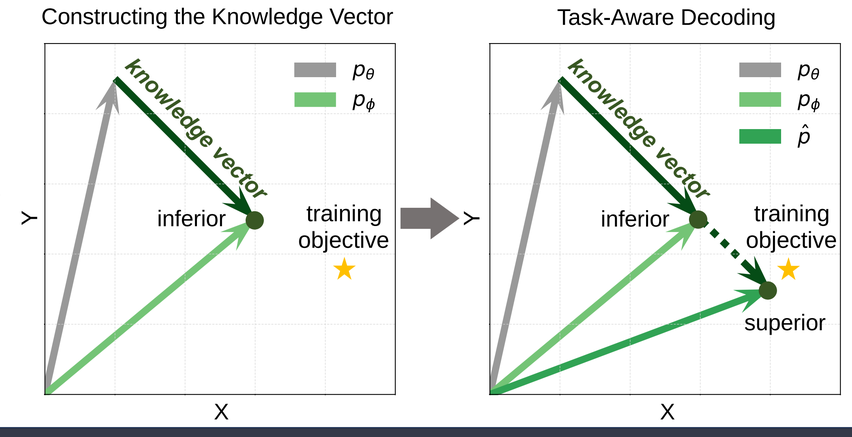

Experiments

- **Results on multiple-choice and CBQA tasks:**

Model	Method	Multiple Choices			CBQA
		MC1	MC2	MC3	True*Info
GPT-J-6b	LoRA	30.6	51.3	25.6	35.7
	+*TaD*	**33.0**	**52.5**	**27.1**	**37.0**
	AdapterP	34.9	54.3	28.0	51.5
	+*TaD*	**38.2**	**55.5**	**29.5**	**51.7**
	AdapterH	36.4	55.0	28.5	53.0
	+*TaD*	**38.3**	**55.8**	**28.7**	**55.3**
	Parallel	34.3	54.0	27.7	47.2
	+*TaD*	**37.5**	**55.1**	**28.9**	**47.4**
BLOOMz-7b	LoRA	30.8	51.4	25.7	17.4
	+*TaD*	**32.8**	**52.3**	**27.2**	**17.5**
	AdapterP	35.3	53.8	**28.5**	20.6
	+*TaD*	**35.7**	**54.8**	28.4	**20.7**
	AdapterH	36.8	54.5	28.9	50.3
	+*TaD*	**37.9**	**55.2**	**29.2**	**50.8**
	Parallel	34.5	53.6	28.2	21.8
	+*TaD*	**36.5**	**54.4**	**28.5**	**22.7**

Model	Method	Multiple Choices			CBQA
		MC1	MC2	MC3	True*Info
LLaMa-7b	LoRA	32.9	55.0	28.5	49.1
	+*TaD*	**34.2**	**55.7**	**29.0**	**51.2**
	AdapterP	38.1	57.4	30.8	61.4
	+*TaD*	**40.6**	**58.5**	**32.1**	**61.8**
	AdapterH	37.8	57.6	30.3	60.3
	+*TaD*	**39.8**	**59.0**	**32.0**	**61.0**
	Parallel	37.0	56.3	29.5	54.3
	+*TaD*	**39.5**	**57.0**	**30.4**	**55.2**
LLaMa-13b	LoRA	33.4	55.7	29.0	54.1
	+*TaD*	**35.1**	**56.7**	**29.7**	**54.7**
	AdapterP	40.6	58.8	32.4	58.6
	+*TaD*	**42.6**	**60.0**	**33.1**	**60.0**
	AdapterH	38.2	57.0	30.4	61.8
	+*TaD*	**39.5**	**57.8**	**31.2**	**63.3**
	Parallel	39.8	58.2	31.7	60.0
	+*TaD*	**42.0**	**60.2**	**33.8**	**61.6**

- **Results on reasoning tasks:**

Model	Method	Math Reasoning		CS Reasoning	
		GSM8K	MultiArith	BoolQ	PIQA
GPT-J-6b	LoRA	21.9	92.5	61.8	63.4
	+*TaD*	**22.8**	**94.2**	**62.7**	**64.6**
	AdapterP	19.0	92.2	63.9	71.0
	+*TaD*	**19.5**	**92.5**	**64.2**	**71.2**
BLOOMz-7b	LoRA	18.9	91.7	66.8	73.6
	+*TaD*	**19.3**	**94.2**	**66.9**	**73.9**
	AdapterP	16.3	90.7	66.2	74.4
	+*TaD*	**17.1**	**93.0**	66.2	**75.0**
LLaMa-7b	LoRA	26.6	90.5	68.7	78.9
	+*TaD*	**27.7**	**91.0**	**69.3**	**79.5**
	AdapterP	31.5	93.5	65.4	76.3
	+*TaD*	**32.0**	**93.7**	**66.3**	76.3
LLaMa-13b	LoRA	35.9	91.5	70.1	82.5
	+*TaD*	**38.1**	**92.0**	**70.8**	**83.1**
	AdapterP	36.8	91.5	69.4	78.1
	+*TaD*	**37.5**	**94.0**	69.4	**79.2**

- **Ablation study of the knowledge vector:**

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
7b	/	10.8/37.5
7b*	/	26.6/90.5
13b	/	16.7/53.2
13b*	/	35.9/91.5

(a) Comparison results on pre-trained and fine-tuned models.

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
7b*	/	26.6/90.5
	7b → 7b*	27.7/91.0
13b*	/	35.9/91.5
	13b →13b*	38.1/92.0

(b) *TaD*'s effectiveness on the fine-tuned models.

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
7b*	/	26.6/90.5
	7b*→ 7b	23.7/79.0

(c) The effect of the opposite direction of the proposed *knowledge vector* (from the fine-tuned to the pre-trained model).

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
13b	/	16.7/53.2
	7b →13b	17.2/51.8
13b*	/	35.9/91.5
	7b*→13b*	36.2/91.8

(d) The effect of the direction of the model size difference (from the smaller to the larger model).

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
7b	/	10.8/37.5
	7b → 7b*	11.9/38.1
	7b →13b	11.2/37.6

(e) Comparison results on the direction of the knowledge and model size difference.

$\mathcal{M}$	$p_\mathcal{S} \to p_\mathcal{E}$	G / M
13b*	/	35.9/91.5
	7b*→13b*	36.2/91.8
	13b →13b*	38.1/92.0
	7b →13b*	38.2/92.0

(f) The cumulative effect of the direction of the knowledge and model size difference.

- **Comparison with other decoding strategies:**

Model	Method	Multiple Choices			Math Reasoning	
		MC1	MC2	MC3	GSM8K	MultiArith
LLaMa-7b	LoRA	32.9	55.0	28.5	26.6	90.5
	+*DoLa*	31.6	48.6	22.7	26.6	89.7
	+*TaD*	**34.2**	**55.7**	**29.0**	**27.7**	**91.0**
	AdapterP	38.1	57.4	30.8	31.5	93.5
	+*DoLa*	39.7	54.9	25.5	31.5	93.3
	+*TaD*	**40.6**	**58.5**	**32.1**	**32.0**	**93.7**
LLaMa-13b	LoRA	33.4	55.7	29.0	35.9	91.5
	+*CD*	**36.2**	55.4	26.5	19.0	70.3
	+*DoLa*	34.9	51.2	24.8	38.0	**94.2**
	+*TaD*	35.1	**56.7**	**29.7**	**38.1**	92.0
	AdapterP	40.6	58.8	32.4	36.8	91.5
	+*CD*	41.1	56.0	26.2	17.8	72.5
	+*DoLa*	41.3	56.5	27.5	35.9	93.5
	+*TaD*	**42.6**	**60.0**	**33.1**	**37.5**	**94.0**

- **Integrated with different basic decoding strategies:**

Model	Method	G / M
LLaMa-7b	Greedy	26.6/90.5
	+*TaD*	**27.7/91.0**
	Beam-4	30.5/91.3
	+*TaD*	**30.9/91.8**
	Top-p	26.7/90.7
	+*TaD*	**27.4/91.3**
	Top-k	27.0/90.3
	+*TaD*	**27.7/91.6**

Model	Method	G / M
LLaMa-13b	Greedy	35.9/91.5
	+*TaD*	**38.1/92.0**
	Beam-4	43.6/93.3
	+*TaD*	**43.7/94.3**
	Top-p	36.7/91.7
	+*TaD*	**37.1/93.0**
	Top-k	36.8/91.7
	+*TaD*	**37.2/93.0**

- **Different ratios of training data and the selection of step:**

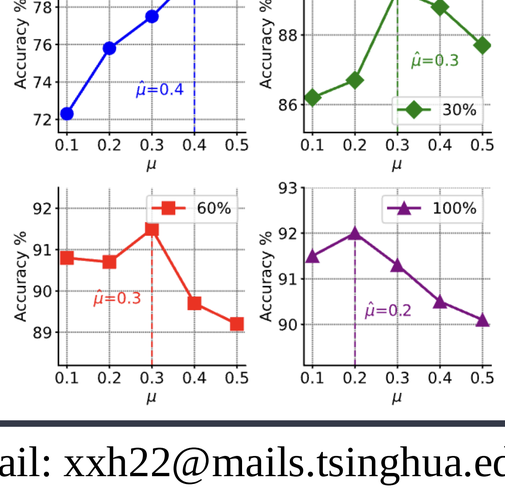

School of Software, BNRist, Tsinghua University, China
Mail: xxh22@mails.tsinghua.edu.cn